

API Exploration Process

1. Initial Observations:

- I started with the v1 version of the API. Initially, I noticed that it returned names starting with the query string and was limited to only 10 results per request.
- Since the URL contained "v1," I assumed there might be multiple versions. Upon further investigation, I found v2 and v3.
- I carried out an Nmap scan to pinpoint the open ports. I discovered that both port 22 (SSH) and port 8000 (API) were open. On **port 8000**, I used `nmap -A` to discover the **backend server** was running **Uvicorn** (`nmap -sV -A -p 8000 {url}`). On port 22, it required a private key, so I stopped.
- Used **SQLMap** to check for **SQL injection vulnerabilities**, assuming SQL might be used in the backend, but found **no vulnerabilities**.
- I hypothesized that a backend function was dynamically generating results based on queries.
- We attempted to reconstruct the backend function using reverse engineering techniques, but found no conclusive patterns. If a function was involved, it was either too complex or not easily reversible.

2. Exploring Endpoints:

- The documentation suggested trying different endpoints, so I used **dirb** to enumerate them. This led me to discover the **/hint**, **/help**, and **/solution** endpoints.

3. Requesting a Hint:

- Sending a **GET** request to the hint endpoint returned a message indicating that I needed a "good way to ask for a hint."
- I experimented with different types of requests and found that a **POST** request successfully provided a hint.

4. Finding the Query Parameter for Increasing Results:

- The hint suggested looking for a query parameter that would increase the number of names returned per request.
- I initially tried generic parameters like `count` and `cnt`, but they did not work.
- After seeking suggestions from ChatGPT, I shifted my focus to parameters related to requests. Eventually, I discovered that `max_requests` was the correct parameter—this was a tough guess.

5. Extracting Names from Different API Versions:

- I carefully analyzed the responses and identified the characters involved in the names.

`v1_char_set: a to z`

`v2_char_set: 0 to 9 and a to z`

`v3_char_set: space, +, -, ., 0 to 9 and a to z`

- Querying with a single character did not return the full list due to the **50-result limit**.
- To bypass this, I generated all **two-character combinations** and queried the API to retrieve names.

6. Handling the Result Limit Recursively:

- Some two-character queries returned more than 50 results. In such cases, I recursively queried using **three-character combinations**.
- If three-character combinations still exceeded the limit, I extended to **four-character combinations** where the first three characters remained fixed.

7. Adapting to API Version Differences:

- v1 and v2 returned names ranging from **2 to 10 characters**, whereas v3 included names **from 1 to 10 characters**.
- For v3, I started with **single-character queries** and followed the same recursive approach to extract all names.

8. Rate Limiting Challenges:

- Each API version imposed rate limits:
 - **v1:** 100 requests
 - **v2:** 50 requests
 - **v3:** 80 requests
- To handle this, I **implemented request throttling** using Python's `time` module to ensure I stayed within the allowed limits.
- I explored IP rotation techniques (ProxyChains) to test API rate limits, but streamlined the process for efficiency.

9. Verification and comparison with the solution endpoint:

- After extracting the data, I verified my results using the **solution endpoint**, which returned an encoded response. I encoded it using the Base64 method, yielding the following result:

`v1_size: 18632`

`v2_size: 13730`

`v3_size: 12517`

`v1_char_set: a-z`

`v2_char_set: a-z0-9`

`v3_char_set: a-z0-9+-.`

`v1_query_params: query=ado&max_results=10`

`max_results_range=1-50`

`v2_query_params: query=ado&max_results=12`

`max_results_range=1-75`

`v3_query_params: query=ado&max_results=15`

`max_results_range=1-100`

- My extracted results were:

- **v1_size:** 18,609
- **v2_size:** 13,701
- **v3_size:** 11,275
- There was a slight difference between my extracted values and the official values, likely due to missing edge cases or unseen constraints.

10. Final Thoughts:

- The process involved **exploring unknown endpoints**, **deciphering undocumented parameters**, and **handling rate limits effectively**.
- It was a challenging yet insightful experience in **API enumeration**, **request handling**, and **adaptive querying strategies**.
- The final verification step helped in identifying discrepancies and areas for further optimization.

Findings of the API

The API is a **RESTful autocomplete service** that takes a few characters as a query and returns a list of names that start with the given input.

Available Endpoints

1. Help Endpoint

- **POST** <http://35.200.185.69:8000/help>
- It provides guidance on how to proceed with the API exploration.
- Suggests trying the **/hint** endpoint.
- Mentions that **rate limiting** is enforced on the IP address.

2. Hint Endpoint

- `POST http://35.200.185.69:8000/hint`
- Provides a hint when accessed with a **POST** request.

3. Solution Endpoint

- `GET http://35.200.185.69:8000/solution`
- Returns encoded statistics about the solution.

4. Version 1 (v1) Autocomplete Endpoint

- `GET http://35.200.185.69:8000/v1/autocomplete?query=a`
- By default, it returns **10 results**, but this can be adjusted using the `max_results` parameter within the range of **1 to 50**.

5. Version 2 (v2) Autocomplete Endpoint

- `GET http://35.200.185.69:8000/v2/autocomplete?query=a`
- By default, it returns **12 results**, but the `max_results` parameter allows modification within the range of **1 to 75**.

6. Version 3 (v3) Autocomplete Endpoint

- `GET http://35.200.185.69:8000/v3/autocomplete?query=a`
- By default, it returns **15 results**, but the `max_results` parameter can be used to modify the count within the range of **1 to 100**.